

Using HPC systems for Python-based AI/ML tasks

UM Fall HPC/Cloud Workshop

Grigory Shamov, October 16, 2025



**University
of Manitoba**

Why the talk is about Python AI/ML on HPC?

- **Everyone (almost) is using Python for AI / ML**
- **HPC has high performance hardware and software**
 - NVIDIA (and AMD) GPUs
 - High bandwidth, Low-latency Interconnect (NVidia Infiniband)
 - Scalable, parallel File Systems (VAST, CEPH, DDN Infinia) or local SSD NVMe
- **HPC community has a data centre infrastructure**
 - Power; high efficiency cooling
- **Many AI shops do use same HPC technology**
 - **SLURM, Linux, etc.**
 - Some do use Kubernetes and containerized workflows



This talk is about using our HPC infra

- HPC presents a few issues for running AI workflows:
 - How to maintain the (largely Python-based) software on HPC?
 - Need to be able to get software and AI Models running
 - Good to have some ability to interactive workflows for debugging
 - How to adapt it to the SLURM based, batch workflow
 - Need to match HPC resources with AI workflow efficiently
- SaaS competition! (Google Colab, for example).

Artificial Intelligence and Machine Learning revolution

<https://static.googleusercontent.com/media/research.google.com/en/pubs/archive/35179.pdf> "The Unreasonable Efficiency of Data"

<http://www.incompleteideas.net/IncIdeas/BitterLesson.html> "The Bitter Lesson"


<https://dl.acm.org/doi/10.5555/3295222.3295349> "Attention is all you need"

<https://blogs.nvidia.com/blog/what-are-foundation-models/> Nvidia on foundational models.



What are those AI and “models”

- Layers of numerical weights, and metadata about them
- Specific to the model’s software and hardware (precision)
 - (“transformers” would have attention layer, etc.)



- **Training/tuning** optimizes the weights of all or some layers
- **Inference** converts a prompt into a numerical representation , passes through model layers and generates outputs

<https://poloclub.github.io/transformer-explainer/>

umanitoba.ca



University
of Manitoba

Working with AI software and “models”

- Python packages : matrix math and optimization (*torch*, *arrow*, *JAX*, ...), ML models (*transformers*, *diffusers*, ...) , training (*accelerate*, *LoRa*, etc.)
 - Need to manage software dependencies (modules, pip, virtualenv)
 - Must use optimized libraries from the host (CUDA, MPI, BLAS/LAPACK, Arrow) → modules
 - Common tools like *conda* are often unwelcome in HPC
- The AI Models themselves (layers, weights) need to be obtained
- A popular community hub Huggingface : <https://huggingface.co/>
- <https://docs.alliancecan.ca/wiki/Huggingface>

Typical tasks for AI / ML on HPC

- Training new models : optimize the weights
- Tuning existing models on additional datasets (LoRA, etc., but more economical: selected layers)
- Inference: getting predictions/generalizations out of existing models
 - Tokenize the “prompt”,
 - Feed it through the layers
 - Convert the representation back to the “output”
- For this Session we will try three examples: Interactive inference for text, text to graphics and tuning/training a graphic model in a batch job.
- Using Grex (mainly) or Magic Castle: need GPUs!

Managing Python dependencies with pip

- Dependencies: Pip handles (mostly) Python dependencies, relying on OS for non-Python ones
 - a. Unlike conda, uv etc, that would package all binary dependencies as well.
 - b. HPC folks like to provide their own optimized binary software : load “modules”!
- Repositories: pip fetches packages from <https://pypi.org/> or a local repository.
 - a. On CCEnv, each and every package must be repackaged to their “wheelhouse”.
 - b. On local software stack SBEnv, manylinux wheels can be used from pypi.org
- Installation destination: in particular when invoked in a Jupyter notebook cell.
 - a. user’s home directory `$HOME/.local/{bin,lib,share}` . (Grex)
 - b. Throw-away virtual environment under `$SLURM_TMPDIR` (Alliance HPC , MagicCastle)
 - c. Explicit virtualenv (create, activate and install) : best method. Needs a Jupyter kernel added.

1. Project one: text generation using Qwen3

- We will generate text based on a prompt, with reasoning using a small LLM
- <https://huggingface.co/Qwen/Qwen3-0.6B>
- The code has three general stages
 1. But first, use pip to install packages in Jupyter.
 - 1.1. Don't forget to load Lmod modules (Lmod tab) !
 - 1.2. Monitoring tab for Node and GPU resources.
 2. Import Python packages and see if they load
 3. Set up encoders for the prompt
 4. Run the inference/generation step
 5. Decode the output, print result
- Can be made into a chatbot with tools like Gradio.

2. Project two: inference using SD-1.5

- We will generate images based on a text prompt
- <https://huggingface.co/stable-diffusion-v1-5/stable-diffusion-v1-5>
 - (used to be **runwayml/stable-diffusion-v1-5**)
 - Will use *torch*, *transformers* and *diffusers*
 - Needs a GPU! On Grex or MagicCastle
 - Will use HuggingFace convenient “*pipe*” interface that lets us set a pretrained model and abstract the encoding, decoding steps.
- Let us use same pip environment in Jupyter as for the project 1.

3. Project 3: using LoRa to train SD-1.5

- The SD-1.5 model is unaware of certain things. I want to add some knowledge to it: an ElderScrolls game character, Scamp
 - <https://en.uesp.net/wiki/Lore:Scamp>
 - Will use a handful of images from the fandom wiki (dataset 2)
 - Will use a handful of auto generated images from ChatGPT 4o. (dataset 1)
 -
- Low Rank Adaptation of Large Language Models (LoRa) technique
 - Only adding a few layers, freezing the rest of the model
 - <https://huggingface.co/docs/diffusers/en/training/lora>

```
unet_lora_config = LoraConfig(  
    r=args.rank,  
    lora_alpha=args.rank,  
    init_lora_weights="gaussian",  
    target_modules=["to_k", "to_q", "to_v",  
        "to_out.0"],  
)  
  
unet.add_adapter(unet_lora_config)  
lora_layers = filter(lambda p:  
    p.requires_grad, unet.parameters())
```

3. Project 3: using LoRa to train SD-1.5

- Using Huggingface “*accelerate*” training script from “*diffusers*” repo
 - Get an interactive job on a GPU node
 - Git clone the repo TBD
 - Create a virtualenv (SBEnv on Grex, CCEnv on Jupyter)
 - Install pip packages, including current versions of *diffusers* from git
 - Download / copy the training data
 - Dataset 1 and dataset 2 with “Scamps”.
 - Run a salloc training job, obtaining new LoRa weights.
 - Validate the model, repeat etc.

3. Project 3: using LoRa to train SD-1.5

- Using Jupyter, repeat the exercise 2 but use generation with new layers added.
 - Will use HuggingFace convenient “*pipe*” interface that lets us set a pretrained model and abstract the encoding, decoding steps.
 - Let us use same pip environment in Jupyter as for the project 2.
 - The only change is to add our new extra layers to the model.
 - Try to generate an image about a “Scamp”.

Strategies for larger AI Model training

- Large datasets, take a lot of time on a single GPU
- Large models would not fit into memory of a single GPU
- Distributed parallel training frameworks: .
 - Huggingface *accelerate* ,
 - Microsoft *DeepSpeed* ,
 - etc,
- Data parallelism (split the data across GPUs)
- Pipeline parallelism (distribute model layers)
- Tensor parallelism (domain decomposition)

May rely on collective MPI operations (AllReduce, AllGather etc,)

DRI (Digital Research Infrastructure) for AI

IN MANY WAYS, AI VINDICATES THE “HPC WAY”

- ▶ **AI needs fast interconnects.** We had them, the cloud and the enterprise did not.
 - ▶ Microsoft deployed 40,000 KM of Infiniband, in 2023, built for the HPC market ~1999,
- ▶ **AI needs message passing.** MPI, the message passing interface, was built Open Source in the HPC community, ~1993
 - ▶ Now the standard library for transformer-based generative AI (e.g. ChatGPT, DeepSpeed, OpenAI etc.).
- ▶ **AI needs heterogeneity** – GPUs for general purpose computing – the hardware building block for AI - came out of the HPC world ("GPGPU" ~2004).
- ▶ **AI needs fast, large scale filesystems** – not object stores
- ▶ **AI needs liquid cooling** – even 5 years ago, many datacenter providers were convinced they could just use air, now none are. HPC systems switched to liquid cooling a long time ago.
- ▶ This means AI needs HPC hardware (probably good) and HPC programmers (good if you are one, bad if you need to hire one).

A slide by Dan Stazione, Director of TACC

tamIA, a real AI supercomputer of LavalU



Which system for which workload?



**Digital Research
Alliance** of Canada

System, kind (2025)	# GPU nodes	GPUs per node layout	Interconnect	Storage, PB
TamIA , HPC (Laval)	42 (H100)	4 x NVIDIA HGX H100 SXM	4 x HDR200 Infiniband, non-blocking	?
Vulcan, HPC (UofA)	205 (L40s)	4 x NVIDIA L40s	1x100Gbps Ethernet	5PB
Killarney, HPC (UofT)	168 (L40s) 10(H100)	4 x NVIDIA L40s, 8 x NVIDIA H100 SXM	1x HDR100, 2x HDR200	1.5 PB
HPC systems ?				
Fir, HPC	160	4 x NVidia H100 SXM	1x HDR200 Infiniband, blocking	51PB
Nibi, HPC	36	8 x Nvidia H100 SXM	1x Nokia 200/400G Ethernet	25PB
Trillium, HPC	60	4 x NVidia H100 SXM	1x NDR200/ NDR400 Infiniband	29PB



**University
of Manitoba**